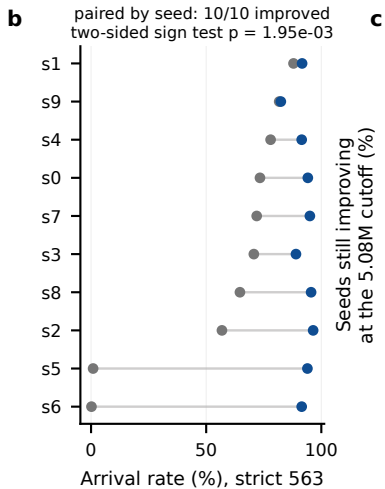
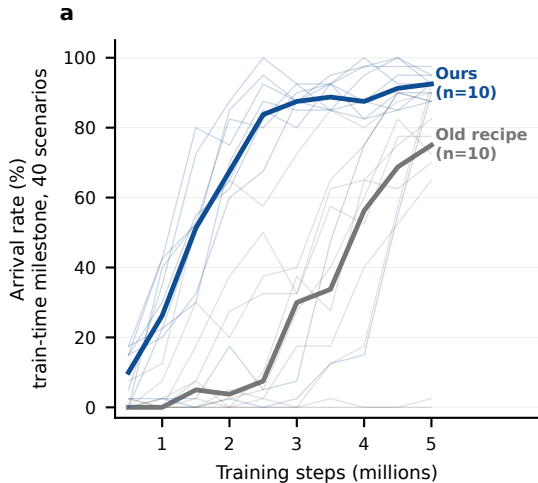


Fig. 1 | Sample efficiency under an identical training budget. a, Learning curves (thin: individual seeds; thick: median). b, Same-seed pairing on the official test set. c, Fraction of non-collapsed seeds whose last three segments still improve by more than two standard errors — the old recipe has not converged at the cutoff, so the comparison is a sample-efficiency claim, not a converged-performance claim. 5.08M-step budget · last checkpoint · official test set (strict 563) · single machine, single pass